

Step 1: Dataset Loading

```
In [1]: import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
In [2]: import pandas as pd
```

```
df = pd.read_csv("sales_data_sample.csv", encoding="latin1")
```

```
In [3]: df.head()
```

```
Out[3]:
```

| | ORDERNUMBER | QUANTITYORDERED | PRICEEACH | ORDERLINENUMBER | SALES | ORDEF |
|---|-------------|-----------------|-----------|-----------------|------------|-------|
| 0 | 10107 | 30 | 95.70 | | 2 2871.00 | 2/24 |
| 1 | 10121 | 34 | 81.35 | | 5 2765.90 | 5/7 |
| 2 | 10134 | 41 | 94.74 | | 2 3884.34 | 7/1 |
| 3 | 10145 | 45 | 83.26 | | 6 3746.70 | 8/25 |
| 4 | 10159 | 49 | 100.00 | | 14 5205.27 | 10/10 |

5 rows × 25 columns



```
In [4]: df.tail()
```

Out[4]:

| | ORDERNUMBER | QUANTITYORDERED | PRICEEACH | ORDERLINENUMBER | SALES | OR |
|-------------|-------------|-----------------|-----------|-----------------|-------|---------|
| 2818 | 10350 | 20 | 100.00 | | 15 | 2244.40 |
| 2819 | 10373 | 29 | 100.00 | | 1 | 3978.51 |
| 2820 | 10386 | 43 | 100.00 | | 4 | 5417.57 |
| 2821 | 10397 | 34 | 62.24 | | 1 | 2116.16 |
| 2822 | 10414 | 47 | 65.52 | | 9 | 3079.44 |

5 rows × 25 columns



In [5]: `df.shape`

Out[5]: (2823, 25)

In [6]: `df.columns`

Out[6]: Index(['ORDERNUMBER', 'QUANTITYORDERED', 'PRICEEACH', 'ORDERLINENUMBER', 'SALES', 'ORDERDATE', 'STATUS', 'QTR_ID', 'MONTH_ID', 'YEAR_ID', 'PRODUCTLINE', 'MSRP', 'PRODUCTCODE', 'CUSTOMERNAME', 'PHONE', 'ADDRESSLINE1', 'ADDRESSLINE2', 'CITY', 'STATE', 'POSTALCODE', 'COUNTRY', 'TERRITORY', 'CONTACTLASTNAME', 'CONTACTFIRSTNAME', 'DEALSIZE'],
dtype='object')

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ORDERNUMBER      2823 non-null    int64  
 1   QUANTITYORDERED 2823 non-null    int64  
 2   PRICEEACH        2823 non-null    float64 
 3   ORDERLINENUMBER 2823 non-null    int64  
 4   SALES            2823 non-null    float64 
 5   ORDERDATE        2823 non-null    object  
 6   STATUS            2823 non-null    object  
 7   QTR_ID           2823 non-null    int64  
 8   MONTH_ID         2823 non-null    int64  
 9   YEAR_ID          2823 non-null    int64  
 10  PRODUCTLINE      2823 non-null    object  
 11  MSRP              2823 non-null    int64  
 12  PRODUCTCODE      2823 non-null    object  
 13  CUSTOMERNAME     2823 non-null    object  
 14  PHONE             2823 non-null    object  
 15  ADDRESSLINE1     2823 non-null    object  
 16  ADDRESSLINE2     302 non-null     object  
 17  CITY              2823 non-null    object  
 18  STATE             1337 non-null    object  
 19  POSTALCODE        2747 non-null    object  
 20  COUNTRY           2823 non-null    object  
 21  TERRITORY         1749 non-null    object  
 22  CONTACTLASTNAME  2823 non-null    object  
 23  CONTACTFIRSTNAME 2823 non-null    object  
 24  DEALSIZE          2823 non-null    object  
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

Step 2: Data Cleaning and Preprocessing

In [8]: `df.isnull().sum()`

```
Out[8]: ORDERNUMBER          0
QUANTITYORDERED         0
PRICEEACH                0
ORDERLINENUMBER          0
SALES                     0
ORDERDATE                 0
STATUS                     0
QTR_ID                    0
MONTH_ID                  0
YEAR_ID                   0
PRODUCTLINE               0
MSRP                      0
PRODUCTCODE               0
CUSTOMERNAME              0
PHONE                     0
ADDRESSLINE1               0
ADDRESSLINE2             2521
CITY                      0
STATE                     1486
POSTALCODE                76
COUNTRY                   0
TERRITORY                 1074
CONTACTLASTNAME           0
CONTACTFIRSTNAME          0
DEALSIZE                  0
dtype: int64
```

```
In [9]: df['ADDRESSLINE2'] = df['ADDRESSLINE2'].fillna("Unknown")
df['POSTALCODE'] = df['POSTALCODE'].fillna("Unknown")
df['TERRITORY'] = df['TERRITORY'].fillna("Unknown")
```

```
In [10]: df.duplicated().sum()
```

```
Out[10]: np.int64(0)
```

```
In [11]: df['ORDERDATE'] = pd.to_datetime(df['ORDERDATE'])
```

```
In [12]: df['COUNTRY'] = df['COUNTRY'].str.title()
df['STATUS'] = df['STATUS'].str.title()
df['DEALSIZE'] = df['DEALSIZE'].str.title()
```

```
In [13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   ORDERNUMBER        2823 non-null    int64  
 1   QUANTITYORDERED   2823 non-null    int64  
 2   PRICEEACH          2823 non-null    float64 
 3   ORDERLINENUMBER   2823 non-null    int64  
 4   SALES              2823 non-null    float64 
 5   ORDERDATE          2823 non-null    datetime64[ns]
 6   STATUS              2823 non-null    object  
 7   QTR_ID              2823 non-null    int64  
 8   MONTH_ID            2823 non-null    int64  
 9   YEAR_ID              2823 non-null    int64  
 10  PRODUCTLINE         2823 non-null    object  
 11  MSRP                2823 non-null    int64  
 12  PRODUCTCODE         2823 non-null    object  
 13  CUSTOMERNAME        2823 non-null    object  
 14  PHONE               2823 non-null    object  
 15  ADDRESSLINE1         2823 non-null    object  
 16  ADDRESSLINE2         2823 non-null    object  
 17  CITY                2823 non-null    object  
 18  STATE               1337 non-null    object  
 19  POSTALCODE           2823 non-null    object  
 20  COUNTRY              2823 non-null    object  
 21  TERRITORY             2823 non-null    object  
 22  CONTACTLASTNAME      2823 non-null    object  
 23  CONTACTFIRSTNAME     2823 non-null    object  
 24  DEALSIZE              2823 non-null    object  
dtypes: datetime64[ns](1), float64(2), int64(7), object(15)
memory usage: 551.5+ KB
```

Step 3: Exploratory Data Analysis (EDA)

In [14]: `df.describe()`

Out[14]:

| | ORDERNUMBER | QUANTITYORDERED | PRICEEACH | ORDERLINENUMBER | SAL |
|--------------|--------------|-----------------|-------------|-----------------|-------------|
| count | 2823.000000 | 2823.000000 | 2823.000000 | 2823.000000 | 2823.000000 |
| mean | 10258.725115 | 35.092809 | 83.658544 | 6.466171 | 3553.8890 |
| min | 10100.000000 | 6.000000 | 26.880000 | 1.000000 | 482.1300 |
| 25% | 10180.000000 | 27.000000 | 68.860000 | 3.000000 | 2203.4300 |
| 50% | 10262.000000 | 35.000000 | 95.700000 | 6.000000 | 3184.8000 |
| 75% | 10333.500000 | 43.000000 | 100.000000 | 9.000000 | 4508.0000 |
| max | 10425.000000 | 97.000000 | 100.000000 | 18.000000 | 14082.8000 |
| std | 92.085478 | 9.741443 | 20.174277 | 4.225841 | 1841.8651 |



In [15]: `df['STATUS'].value_counts()`

Out[15]: STATUS

| | |
|------------|------|
| Shipped | 2617 |
| Cancelled | 60 |
| Resolved | 47 |
| On Hold | 44 |
| In Process | 41 |
| Disputed | 14 |

Name: count, dtype: int64

In [16]: `df.groupby('COUNTRY')['SALES'].sum().sort_values(ascending=False)`

```
Out[16]: COUNTRY
Usa           3627982.83
Spain         1215686.92
France        1110916.52
Australia     630623.10
Uk            478880.46
Italy          374674.31
Finland       329581.91
Norway         307463.70
Singapore      288488.41
Denmark        245637.15
Canada         224078.56
Germany        220472.09
Sweden         210014.21
Austria        202062.53
Japan           188167.81
Switzerland    117713.56
Belgium         108412.62
Philippines    94015.73
Ireland         57756.43
Name: SALES, dtype: float64
```

```
In [17]: df.groupby('PRODUCTLINE')['SALES'].sum().sort_values(ascending=False)
```

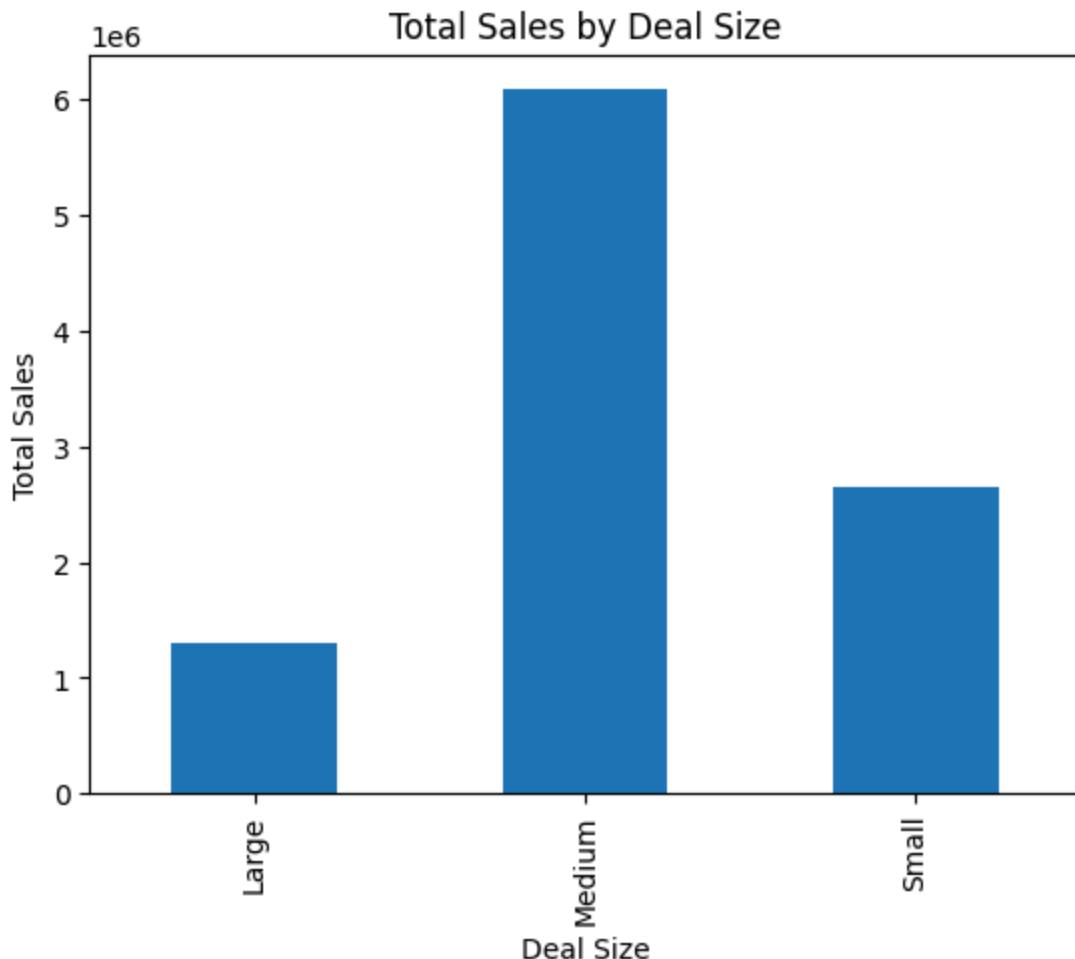
```
Out[17]: PRODUCTLINE
Classic Cars      3919615.66
Vintage Cars      1903150.84
Motorcycles        1166388.34
Trucks and Buses   1127789.84
Planes             975003.57
Ships              714437.13
Trains             226243.47
Name: SALES, dtype: float64
```

```
In [18]: df[['SALES', 'QUANTITYORDERED', 'PRICEEACH']].corr()
```

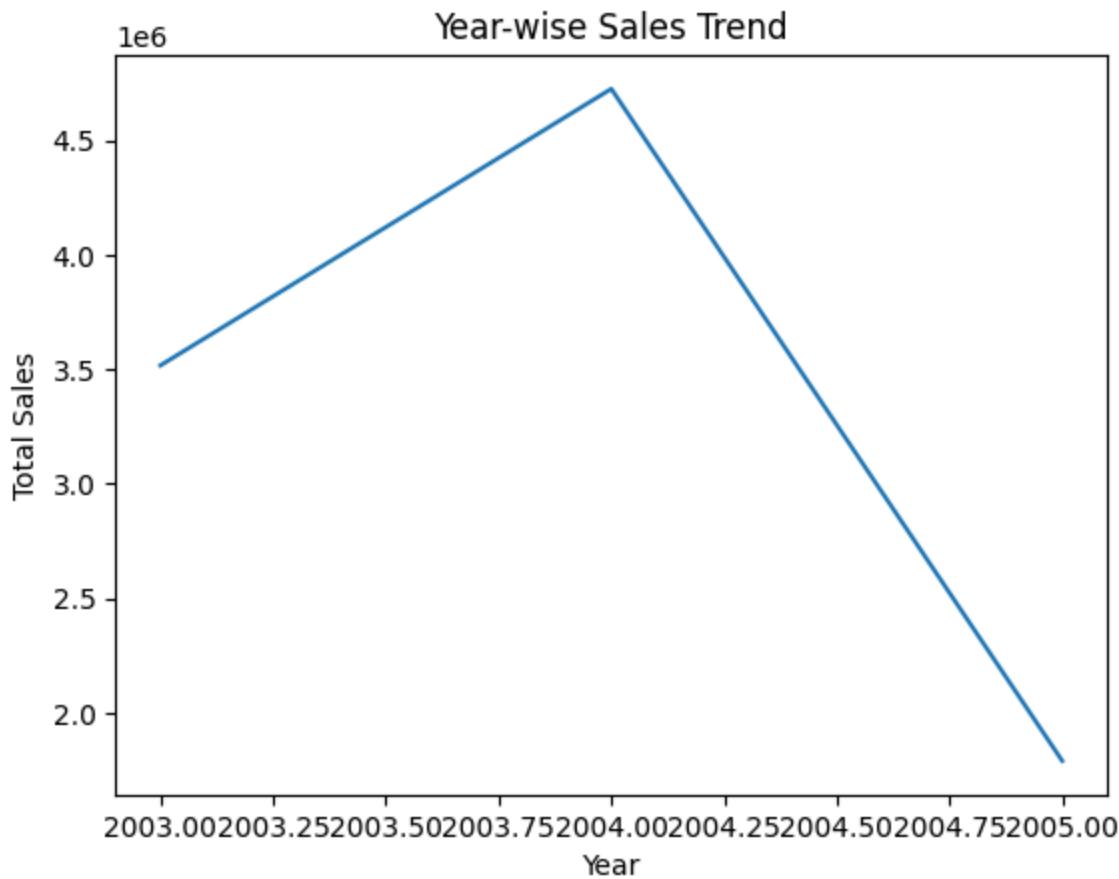
| | SALES | QUANTITYORDERED | PRICEEACH |
|------------------------|----------|-----------------|-----------|
| SALES | 1.000000 | 0.551426 | 0.657841 |
| QUANTITYORDERED | 0.551426 | 1.000000 | 0.005564 |
| PRICEEACH | 0.657841 | 0.005564 | 1.000000 |

Step 4: Data Visualization

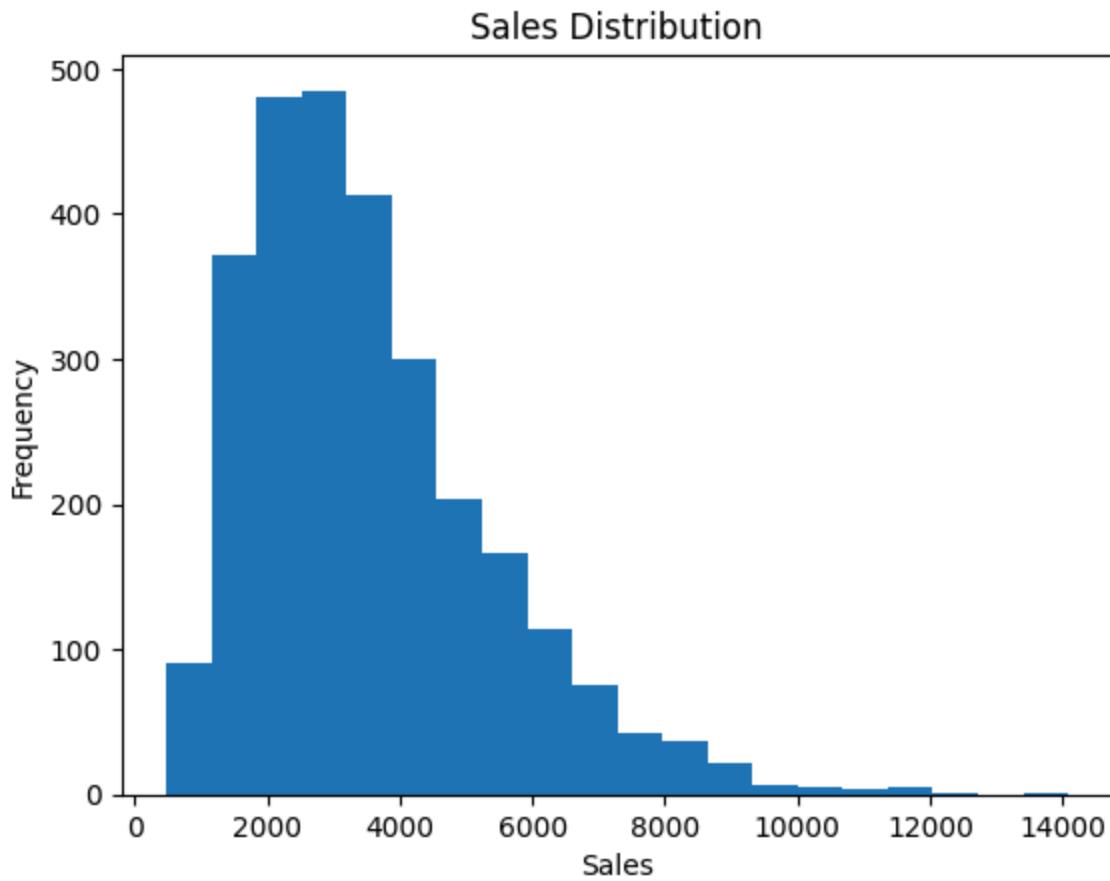
```
In [19]: df.groupby('DEALSIZE')['SALES'].sum().plot(kind='bar')
plt.title("Total Sales by Deal Size")
plt.xlabel("Deal Size")
plt.ylabel("Total Sales")
plt.show()
```



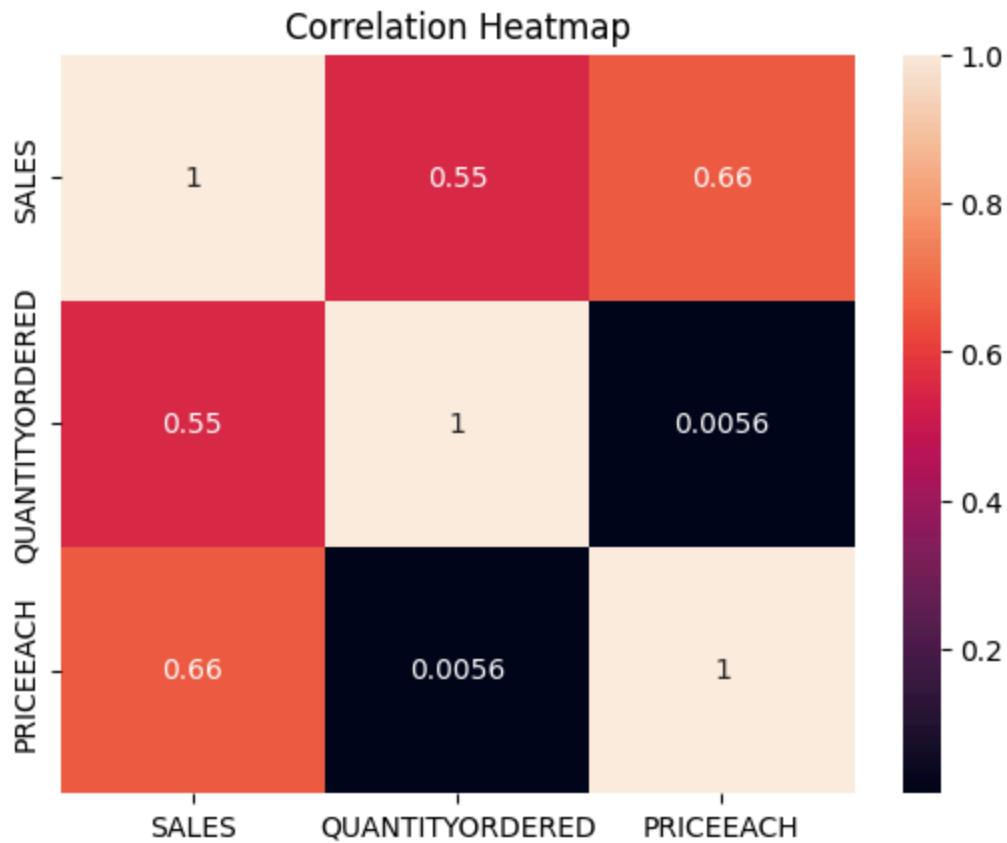
```
In [20]: df.groupby('YEAR_ID')['SALES'].sum().plot(kind='line')
plt.title("Year-wise Sales Trend")
plt.xlabel("Year")
plt.ylabel("Total Sales")
plt.show()
```



```
In [21]: plt.hist(df['SALES'], bins=20)
plt.title("Sales Distribution")
plt.xlabel("Sales")
plt.ylabel("Frequency")
plt.show()
```

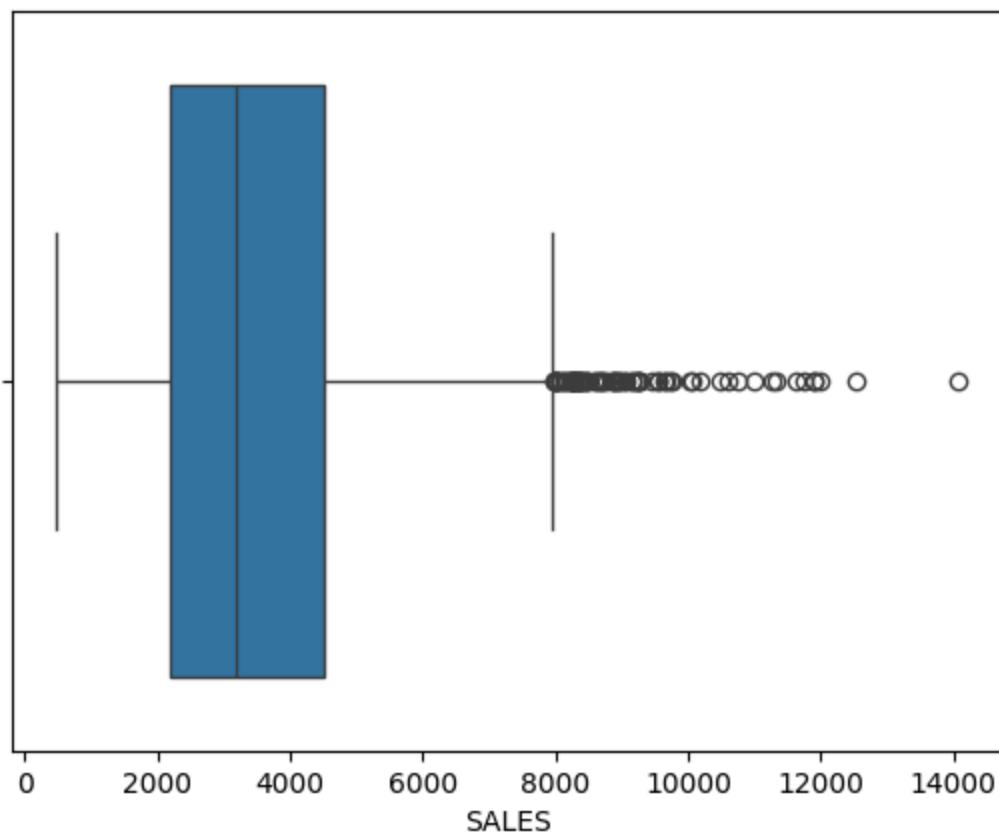


```
In [22]: sns.heatmap(  
    df[['SALES', 'QUANTITYORDERED', 'PRICEEACH']].corr(),  
    annot=True  
)  
plt.title("Correlation Heatmap")  
plt.show()
```



```
In [23]: sns.boxplot(x=df['SALES'])
plt.title("Sales Outliers")
plt.show()
```

Sales Outliers



Step 5: Insights and Interpretation

Sales by Deal Size Large deals generate the highest total sales, followed by medium and small deals, indicating that targeting larger orders significantly boosts revenue.

Year-wise Sales Trend Sales have been increasing year over year, showing a positive growth trend and suggesting that business strategies are effectively driving revenue.

Product Line Performance Quantity ordered shows almost no correlation with price per item, meaning customers do not significantly change order volume based on product price.

Sales Distribution & Outliers Sales values are right-skewed with some extreme high values, indicating that a few large transactions dominate revenue, which could affect forecasting and inventory planning.